

Technologies used:

| | |
|-------------------------------------|---|
| Hadoop (In pseudo-distributed mode) | Hadoop 2.7.7 |
| Java | java version "1.8.0_201" Java(TM) SE Runtime Environment (build 1.8.0_201-b09) Java HotSpot(TM) 64-Bit Server VM (build 25.201-b09, mixed mode) |
| Spark (pyspark shell is used) | Version 2.4.4 |
| Python | Version 3 |
| Flask | Used as framework server |
| Ubuntu | 18.04.1 LTS |

Application Architecture:

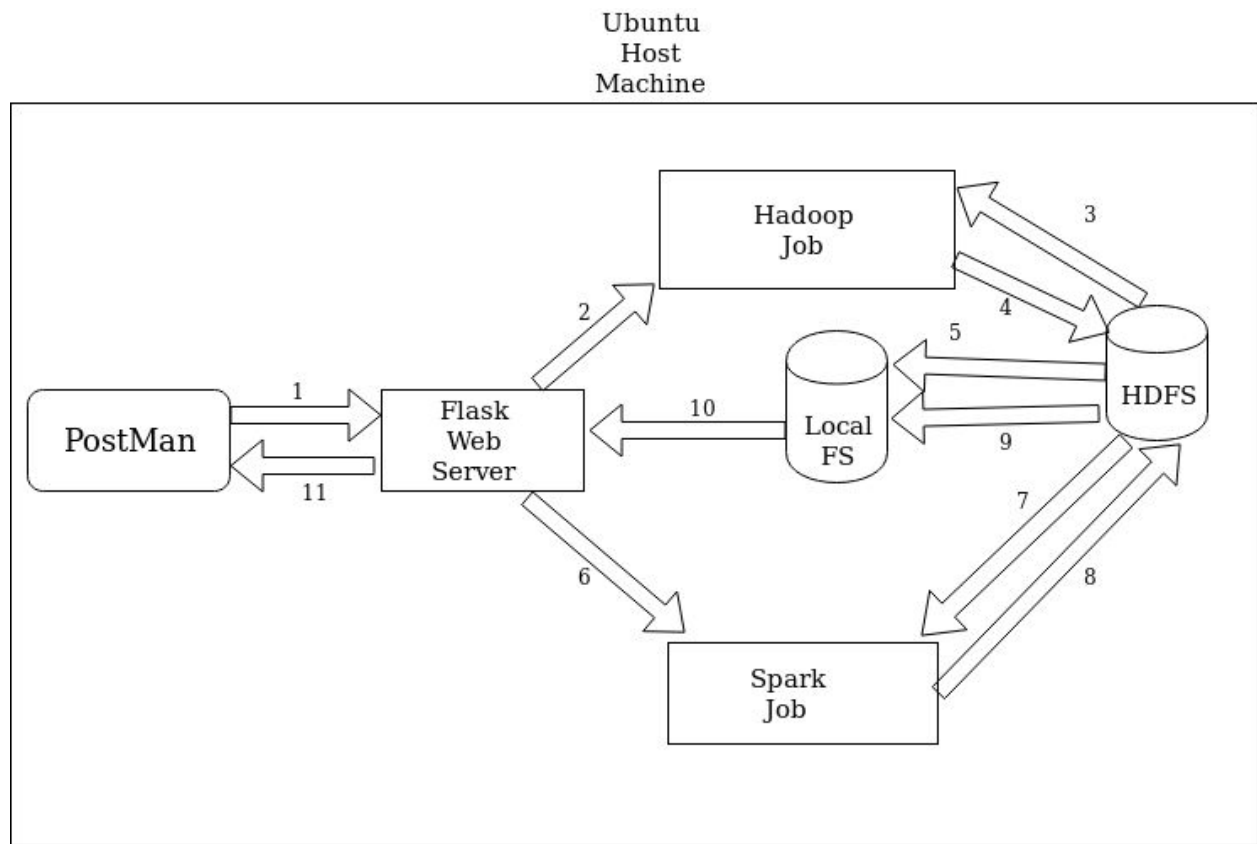


fig. Application Architecture

Description of Application Architecture:

- 1) Client sending request to Flask Web Server
- 2) Flask server forwarding the request to Hadoop Job
- 3) Hadoop Job reading the data from HDFS(Hadoop File System)
- 4) Hadoop Job writing the data to HDFS
- 5) Data gets written to local File System
- 6) Flask server forwarding the request to Spark Job
- 7) Spark Job reading the data from HDFS
- 8) Spark Job writing the data to HDFS
- 9) Data gets written to local File System
- 10) Response being returned to Flask Web Server
- 11) Client getting the final output

The Software Application consists of 3 Modules

1. Restful Web Service Module
2. Hadoop Map-Reduce Module
3. Spark Transformation-Action Module

All the modules are explained in detail below:-

Restful Web Service

Web Service Configuration:

URI : <http://localhost:6006/inputString>

METHOD:POST

BODY:

```
{  
    "query" : " Join Query or Group By Query "  
}
```


Sequence of operations at server

1. Query is provided to the server through the above URI using any Web Service Tool such as Postman
2. After the Query String is received at the **Flask Web Server**, it is **parsed** and classified as **Join Query or Group By Query**
3. Next, all the parameters of the query are extracted and saved at the **designated data structures** at the server and **Hadoop(Function-1)** and **Spark(Function-2)** operations are applied in **sequence**.
4. **Function 1 - Hadoop**
 - a. First server checks if the hadoop **Name Node** and **Data Nodes** are up or not at **http://localhost:9000**
 - b. Next based on the query, server picks up the designated jar and executes it with hadoop along with various **query parameters** extracted previously in step 3.
 - c. After the hadoop job execution, server opens the **YARN Resource Manager** running at **localhost:8088** and scrapes the **latest job ID** of the recently run jobs

←

→

🔒 Not secure | admin44-optiplex-3020:19888/jobhistory/job/job_1569009244950_0149



MapReduce Job job_1569009244950_0149

Application

Job

Overview

Counters

Configuration

Map tasks

Reduce tasks

Tools

Job Name:

Users Zipcodes

User Name:

HDUSER

Queue:

default

State:

SUCCEEDED

Uberized:

false

Submitted:

Mon Sep 23 01:47:29 IST 2019

Started:

Mon Sep 23 01:47:33 IST 2019

Finished:

Mon Sep 23 01:47:42 IST 2019

Elapsed:

8sec

Diagnostics:

Average Map Time:

2sec

Average Shuffle Time:

1sec

Average Merge Time:

0sec

Average Reduce Time:

0sec

ApplicationMaster

| Attempt Number | Start Time | Node |
|----------------|------------------------------|----------------------------|
| 1 | Mon Sep 23 01:47:31 IST 2019 | admin44-OptiPlex-3020-8042 |

| Task Type | Total | Complete | |
|--------------|--------|----------|------------|
| Map | 2 | 2 | |
| Reduce | 1 | 1 | |
| Attempt Type | Failed | Killed | Successful |
| Maps | 0 | 0 | 2 |
| Reduces | 0 | 0 | 1 |

- d. Now as the server has the Job ID for Hadoop Job it opens the URL **http://localhost:19888/jobhistory/job_<job_id>** and scrapes all the **Hadoop Execution Log Parameters** such as Total Elapsed Time, Average Map Time,Average Reducer Time etc..
- e. All the fetched execution parameters are returned

5. Function 2 - Spark

- a. Based on the query, server picks up the designated python script and execute it along with various **query parameters** extracted previously in step 3
 - b. All the **Spark execution log parameters** are fetched based on the OS Timestamp.
6. Now the HDFS file system is read which was previously written by Hadoop and Spark Jobs for the required query output.
 7. The Query output along with execution parameters are returned to the client.

Hadoop Map-Reduce

Query 1:

Join Template:

```
SELECT * FROM <TABLE1> INNER JOIN <TABLE2> ON <CONDITION1>
WHERE <CONDITION2>
```

Example:

```
SELECT * FROM USERS INNER JOIN zipcodes ON zipcode
WHERE userid = 780
```

Description:

For this query we are operating on two tables. One is user table and other is zipcodes table. We are first splitting the user table as key-value pair where **zipcode** is selected as **key** and remaining record is selected as value. For zipcode table, we are selecting **zipcode** as **key** and remaining record is selected as value. After Mapping and Filtering, we are getting Mapper output as:-

i) For user table we are getting records in the form

<zipcode> , <userid, age, gender, occupation, zipcode>

ii) For zipcode table we are getting records in the form

<zipcode>,<zipcode,zipcodetype,city,state>

All the records with the same key (zipcode here) are Reduced to to the same Reducer and joined to form the resultant join response.

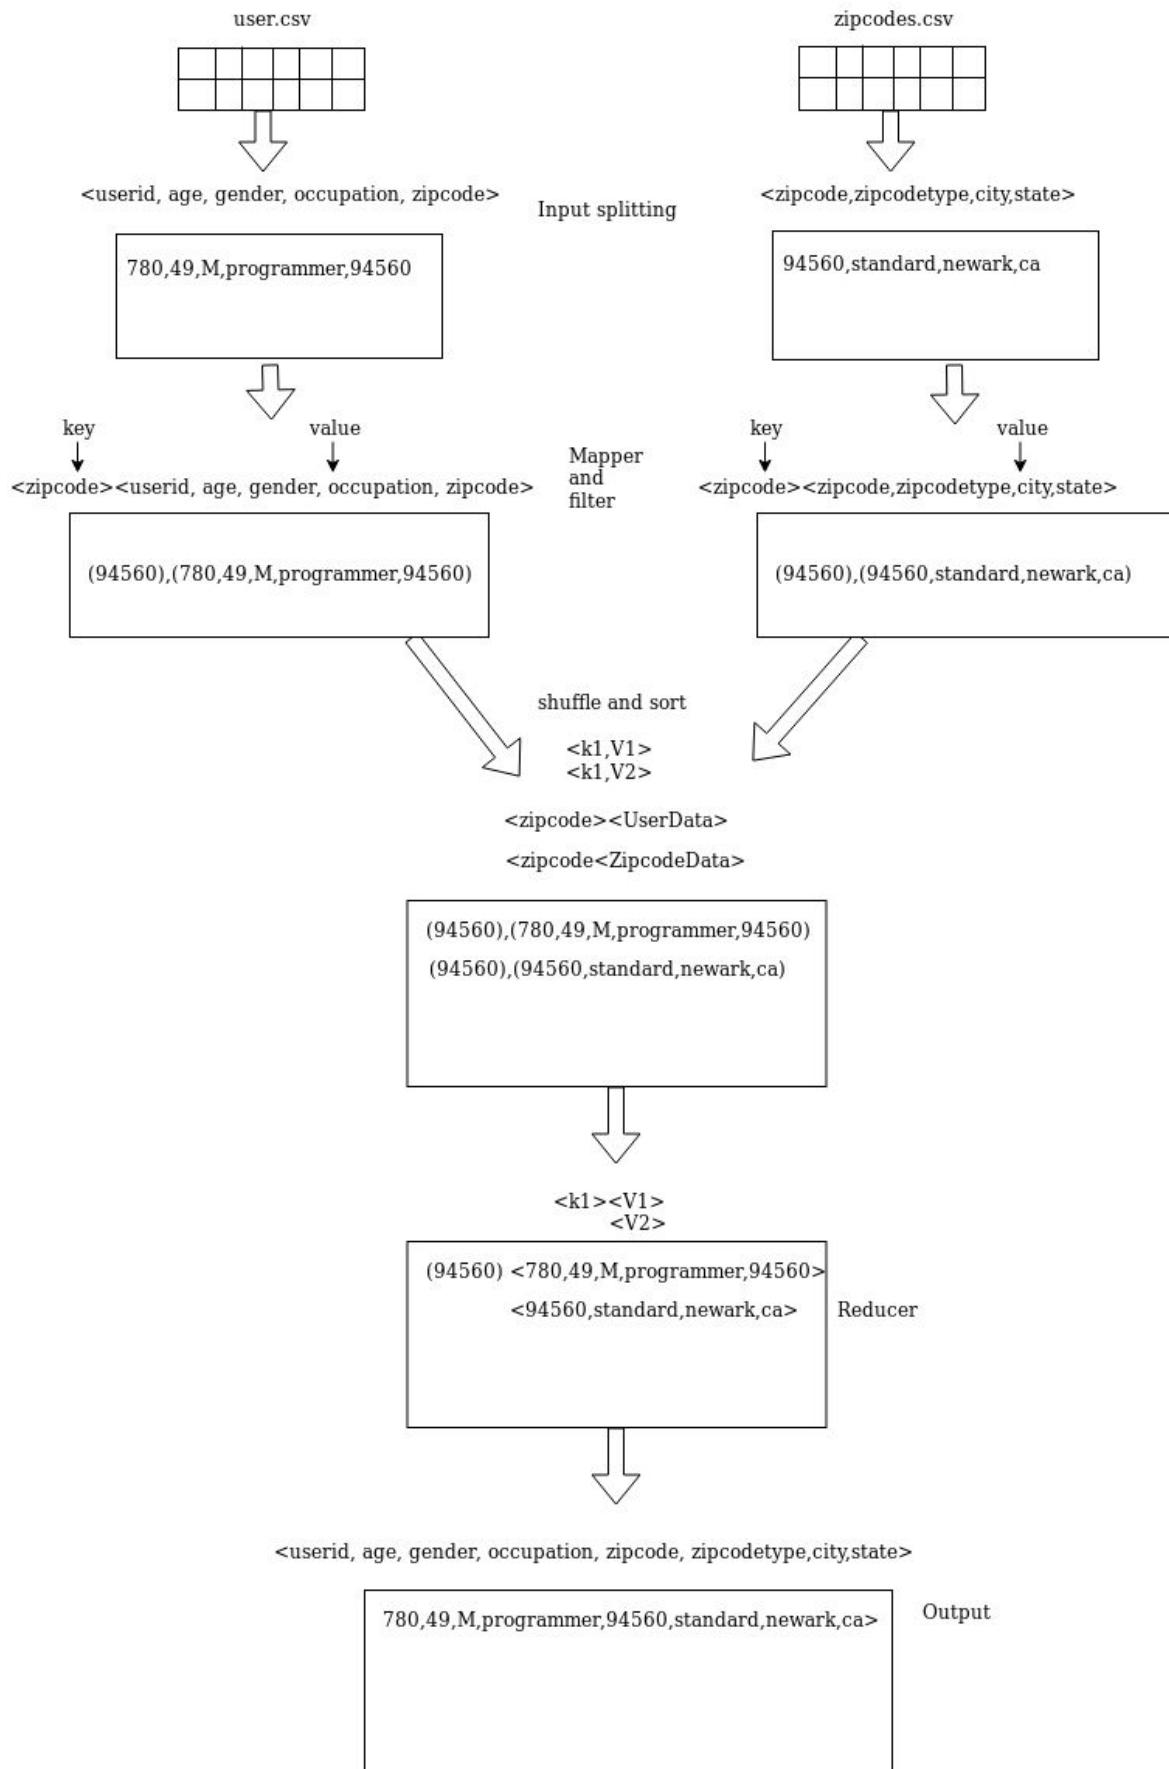


fig. Data Flow Architecture of Map-Reduce: Query 1

Query 2:

Group By Template:

```
SELECT <COLUMNS>, FUNC(COLUMN1)
FROM <TABLE>
GROUP BY <COLUMNS>
HAVING FUNC(COLUMN1)>X
```

Example:

```
SELECT gender, occupation, MAX(age)
FROM users
GROUP BY gender, occupation
HAVING MAX(age) > 45
```

Description:

For this query we are operating on user table. Once we import the input, we are then choosing **gender and occupation as key** and the rest of the record is chosen as **value**. Therefore for user table we are getting records in the form (<**gender, occupation**>, <**userid, age, gender, occupation, zipcode**>). We are combining the records using “shuffle and sort”. Now values are ordered according to key. We are then using reducer. As in the “having” clause, we are selecting only those users who have their age greater than 50. Once reducer acts according to logic supplied we are able to get final output.

users.csv

| | | | | | | |
|--|--|--|--|--|--|--|
| | | | | | | |
| | | | | | | |



<userid, age, gender, occupation, zipcode>

<792,40,M,programmer,12205>
<795,30,M,programmer,8610>
<800,25,M,programmer,55337>
<811,40,F,educator,73013>
<878,50,F,educator,98027>

Input splitting



<key,value>

<(M,programmer),40>
<(M,programmer),30>
<(M,programmer),25>
<(F,educator),40>
<(F,educator),50>

Mapper and filter



<key,value>

<(M,programmer),25>
<(M,programmer),30>
<(M,programmer),40>
<(F,educator),40>
<(F,educator),50>

Shuffle and sort



<key,value>

<M,programmer>,40>
<(F,educator),50>

Reducer



<key,value>

<(F,educator),50>

Output

fig. Data Flow Architecture of Map-Reduce: Query 2

Spark Transformations-Actions

Query 1:

Join Template:

```
SELECT * FROM <TABLE1> INNER JOIN <TABLE2> ON <CONDITION1>
WHERE <CONDITION2>
```

Example:

```
SELECT * FROM users INNER JOIN rating ON userid
WHERE zipcode = 30067
```

Description

Once we start pyspark shell, sparkContext object (sc) is available for us. We could use this object for further processing. As we know Spark framework provides functions for basic transformations and actions, to carry out this query we have used the following **transformations:**

| | |
|----------|--|
| map() | We have used map function to transform user table into key- value pair. First field is selected as key, and the rest of the fields are treated as values. We have done a similar thing for rating table. |
| join() | We are using this transformation function to join user table and rating table. As this function requires two datasets to be in the form of (K,V) pair, we have first converted user table and rating table by using map function and then we are applying join function over them. |
| filter() | To execute where condition in the template, we are using filter function. We are applying |

| | |
|--|--|
| | filter function to the joint table of user and rating (which we obtained in the previous step). To filter function we are providing logic to select those records which have zipcode as “30067”. |
|--|--|

We have used the following **actions** in this task:

| | |
|------------------|--|
| collect() | To print the result of the query, we are using this action provided by spark framework. We are calling collect function on the dataset (RDD in this case) that we need to print. |
| take() | Sometimes we are just verifying outputs of intermediate steps. Now instead of printing entire result, we are using take function provided by spark framework. |
| saveAsTextFile() | At the end we are saving final output of the query in a file. For this we are using this function provided by spark framework. |

First on user table we are applying **Map** function. We are selecting first field which is **userid** as **key** of the table. Map function will **transform** user table as

<userid> ,<userid, age, gender, occupation, zipcode>

Similarly we are selecting first field of rating table i.e. **userid** as **key** of the table. Map function is applied on rating table and table will be transformed to key value pair where key is **userid** and value is (**userid, movieid, rating, timestamp**).

Next we are applying join function on user table and rating table. Condition for join is userid. Now both tables are concatenated. We are then applying filter function on big table. We are selecting users who have zip code as “30067”. We are then using “saveAsTextFile” transformation function to save the RDD in HDFS File System.

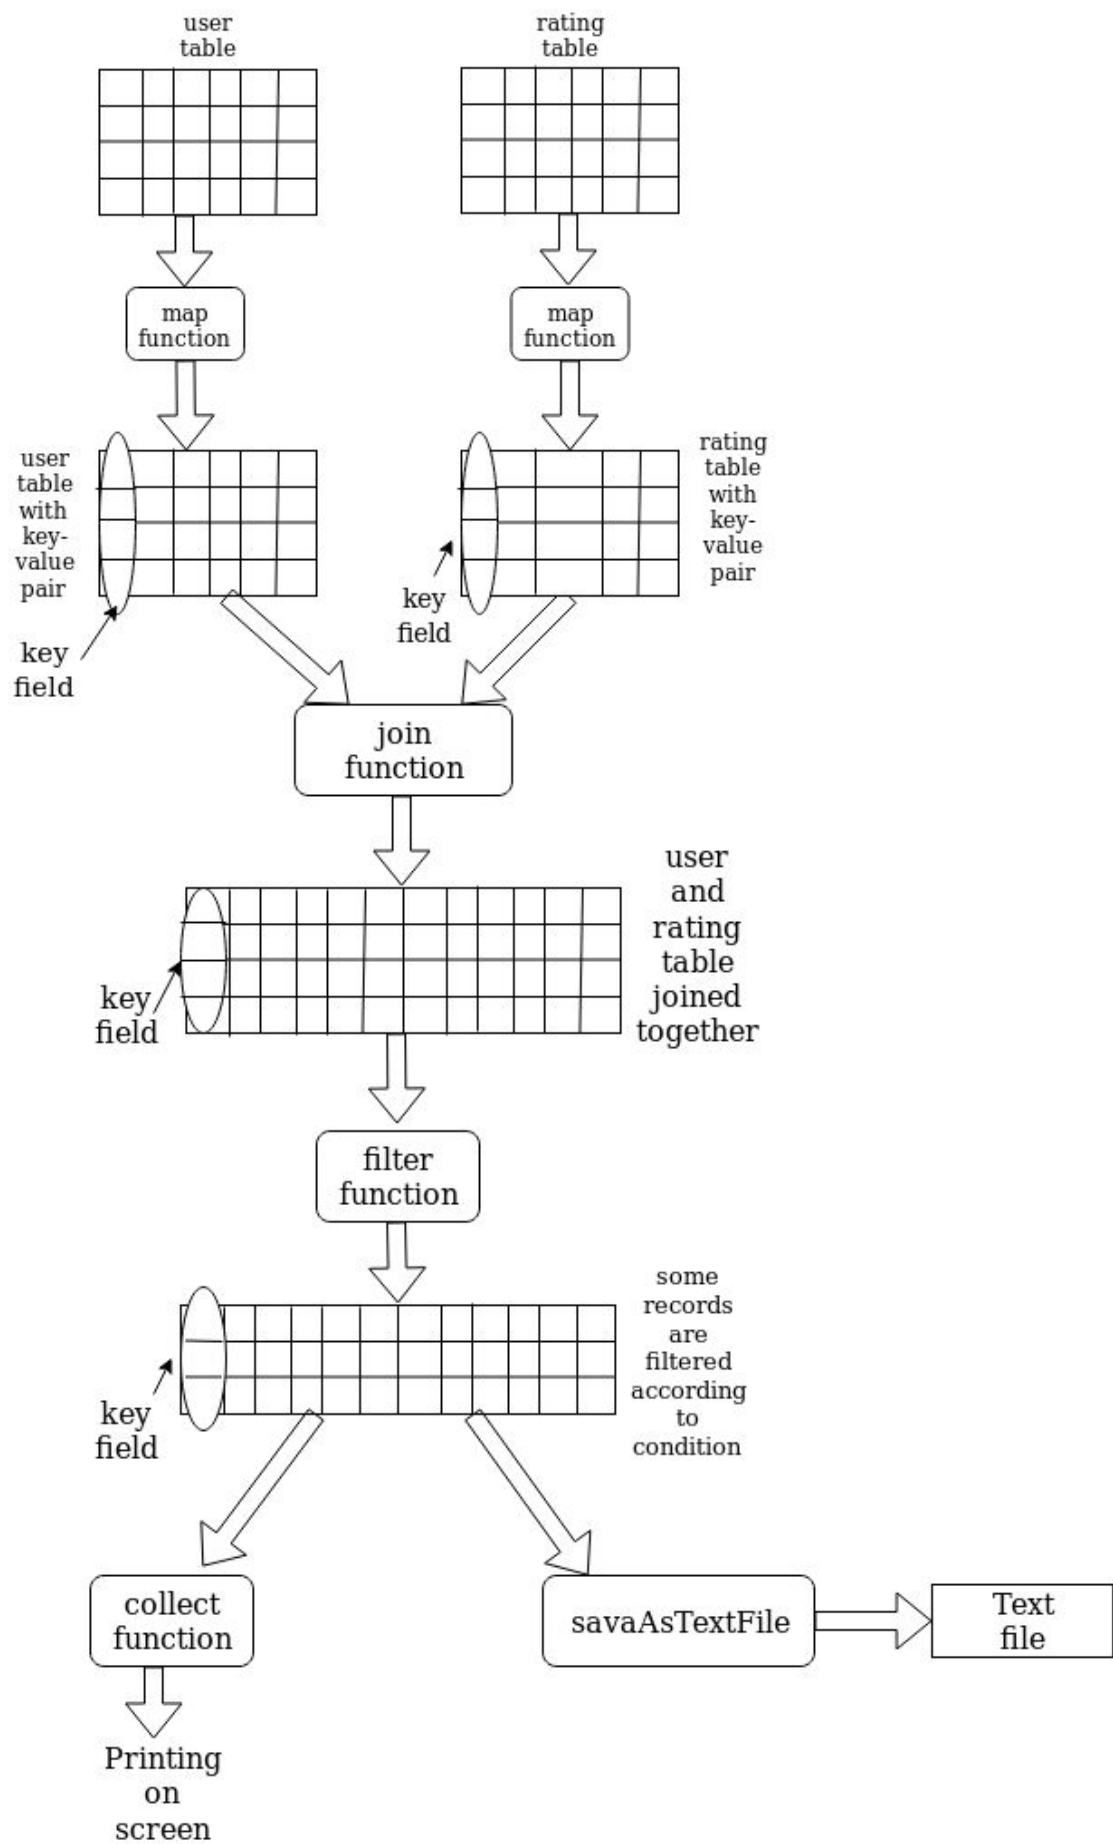


Fig. Data Flow Architecture of Spark: Query 1

Query 2:

Group By Template

```
SELECT <COLUMNS>, FUNC(COLUMN1)
FROM <TABLE>
GROUP BY <COLUMNS>
HAVING FUNC(COLUMN1)>X
```

Example:

```
SELECT occupation,gender, SUM(age)
FROM users
GROUP BY occupation,gender
HAVING SUM(age) > 1000
```

Description

We have used following **transformations** for this task:

| | |
|--------------|--|
| map() | We have used map function is to transform user table into key- value pair. Here we are taking multiple columns as key as per given in group by clause. |
| filter() | As we need to select the records according to condition in having clause, we are using filter transformation for this purpose. |
| groupByKey() | Once we get the columns which are supposed to by keys, we are using groupByKey() transformation, so that data for that particular key could be aggregated. |

We have used the following **actions** for this task:

| | |
|------------------|--|
| collect() | To print the result of the query, we are using this action provided by spark framework. We are calling collect function on the dataset (RDD in this case) that we need to print. |
| take() | Sometimes we are just verifying outputs of intermediate steps. Now instead of printing entire result, we are using take function provided by spark framework. |
| saveAsTextFile() | At the end we are saving final output of the query in a file. For this we are using this function provided by spark framework. |

First on **user** table we are applying **Map** function. As, in group by clause, we have gender and occupation, we are using those fields as key. Map function will transform user table to key value pair as

<gender, occupation>,<userid, age, gender, occupation, zipcode>

We are now applying “groupByKey” function on user table. It returns a dataset of (K, Iterable<V>) pairs.

Now we are iterating over that dataset, finding the sum of ages of users (as they are already grouped by gender and occupation) and mapping it to a new RDD. Now we are applying filter function where condition is mentioned as per the having clause. We are selecting users whose sum of the ages is greater than 1000. We are then using “saveAsTextFile” function to save the output in file.

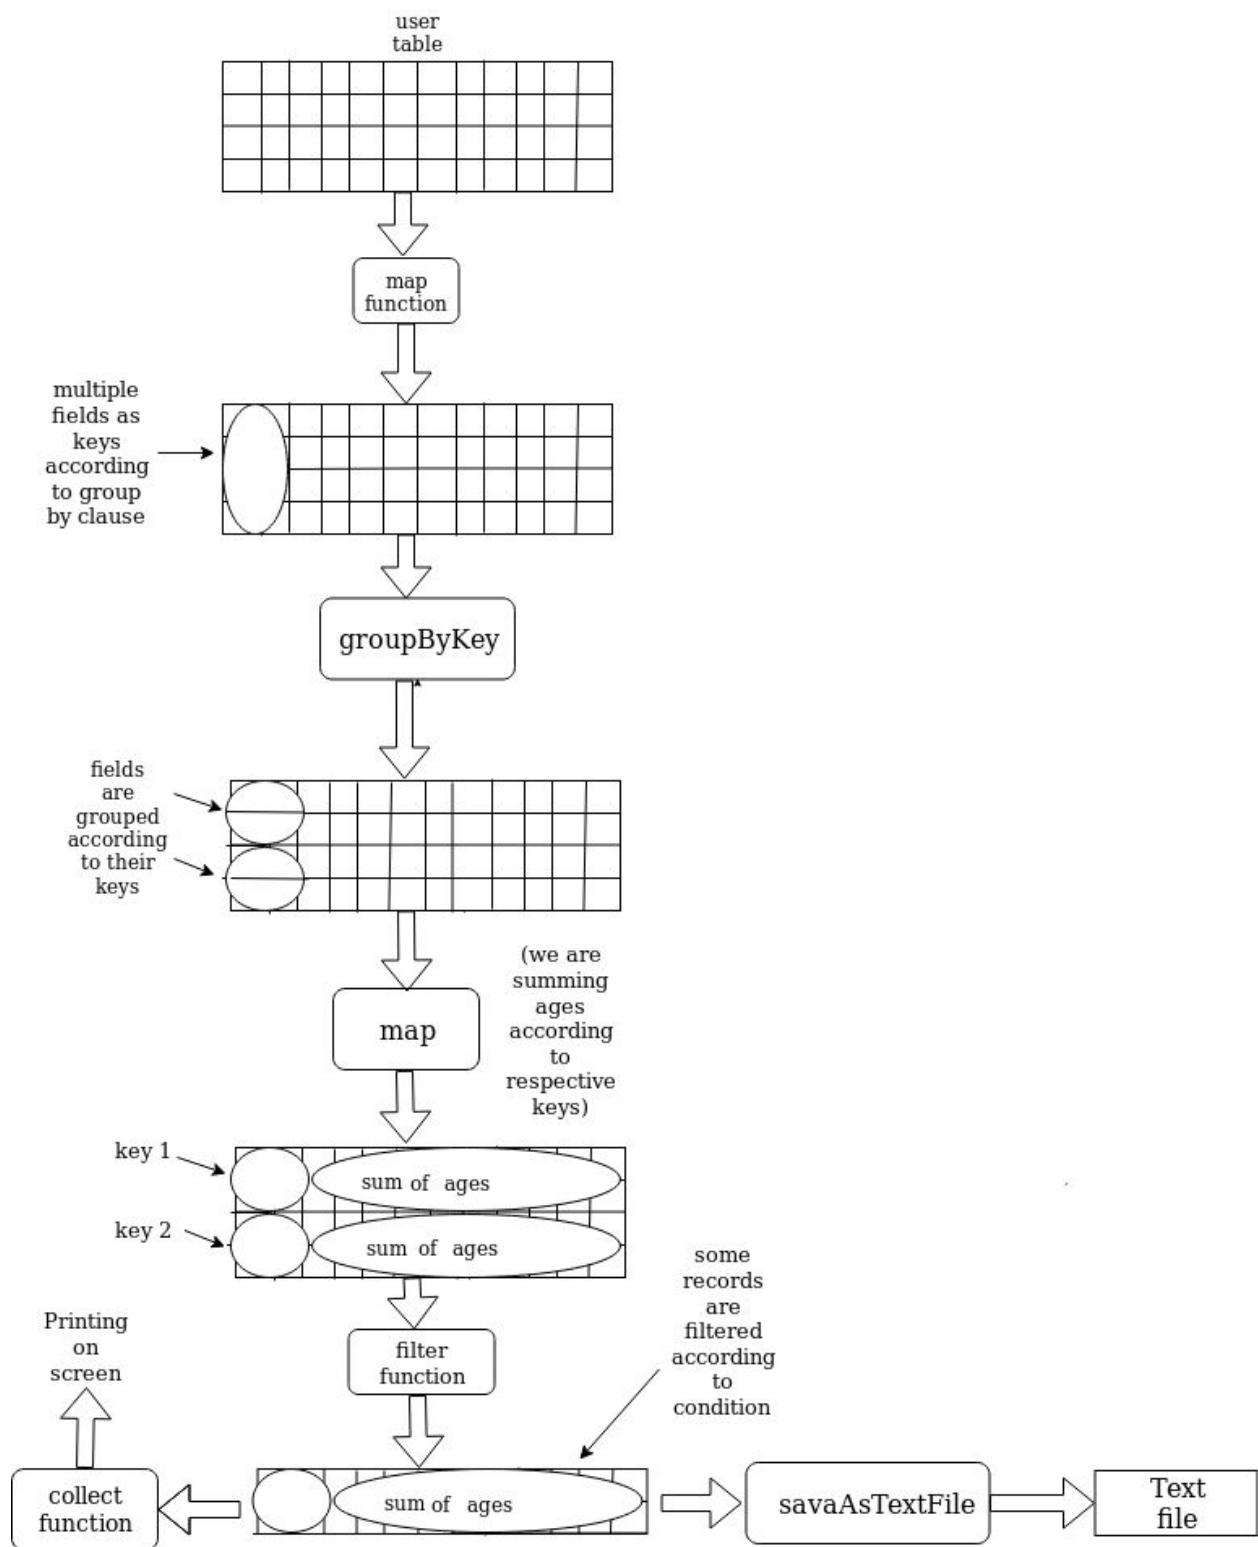


fig. Data Flow Architecture of Spark: Query 2

Outputs

Query 1

The screenshot displays a web interface for managing queries. At the top, there are radio buttons for different content types: none, form-data, x-www-form-urlencoded, raw (selected), binary, and JSON (application/json). Below this, a text area contains a query: `"query": "SELECT * from users inner join zipcodes on zipcode where age = 63"`. The interface then shows the response body, which is a large JSON object. The response includes various performance metrics like `hadoop_average_map_time`, `hadoop_total_elapsed_time`, and `spark_time_execution`. It also contains a `hadoop_query_result` array with three rows of user data and a `map_chain_operations` array detailing the execution steps from Mapper input to Reducer output.

```
1 {
2   "query": "SELECT * from users inner join zipcodes on zipcode where age = 63"
3 }
```

Body Cookies Headers (4) Test Results

Pretty Raw Preview JSON

```
1 {
2   "hadoop_average_map_time": "2sec",
3   "hadoop_average_merge_time": "0sec",
4   "hadoop_average_reduce_time": "0sec",
5   "hadoop_average_shuffle_time": "1sec",
6   "hadoop_query_result": [
7     "777,63,M,programmer,1810,STANDARD,ANDOVER,MA",
8     "364,63,M,engineer,1810,STANDARD,ANDOVER,MA",
9     "858,63,M,educator,9645,MILITARY,FPO,AE"
10  ],
11   "hadoop_total_elapsed_time": "8sec",
12   "map_chain_operations": [
13     "Mapper Input",
14     "Mapper 1 Input: userid,age,gender,occupation,zipcode",
15     "Mapper 1 Output:",
16     "Key: zipcode",
17     "Value: userid,age,gender,occupation,zipcode",
18     "Mapper 2 Input: zipcode,zipcodetype,city,state",
19     "Mapper 2 Output:",
20     "Key:zipcode",
21     "Value:zipcode,zipcodetype,city,state",
22     "Sorting and Shuffling",
23     "Reducer Input: Mapper 1 and Mapper 2 Output belonging to same Key zipcode",
24     "Reducer Output:",
25     "userid,age,gender,occupation,zipcode,zipcodetype,city,state"
26  ],
27   "process_name": "application_1569009244950_0195",
28   "spark_query_result": [
29     "858,63,M,educator,9645,MILITARY,FPO,AE",
30     "364,63,M,engineer,1810,STANDARD,ANDOVER,MA",
31     "777,63,M,programmer,1810,STANDARD,ANDOVER,MA"
32  ],
33   "spark_time_execution": "4.403885364532471"
34 }
```

fig. Input Query Request to Server and Response from Server

Query 2



fig. Input Query Request to Server

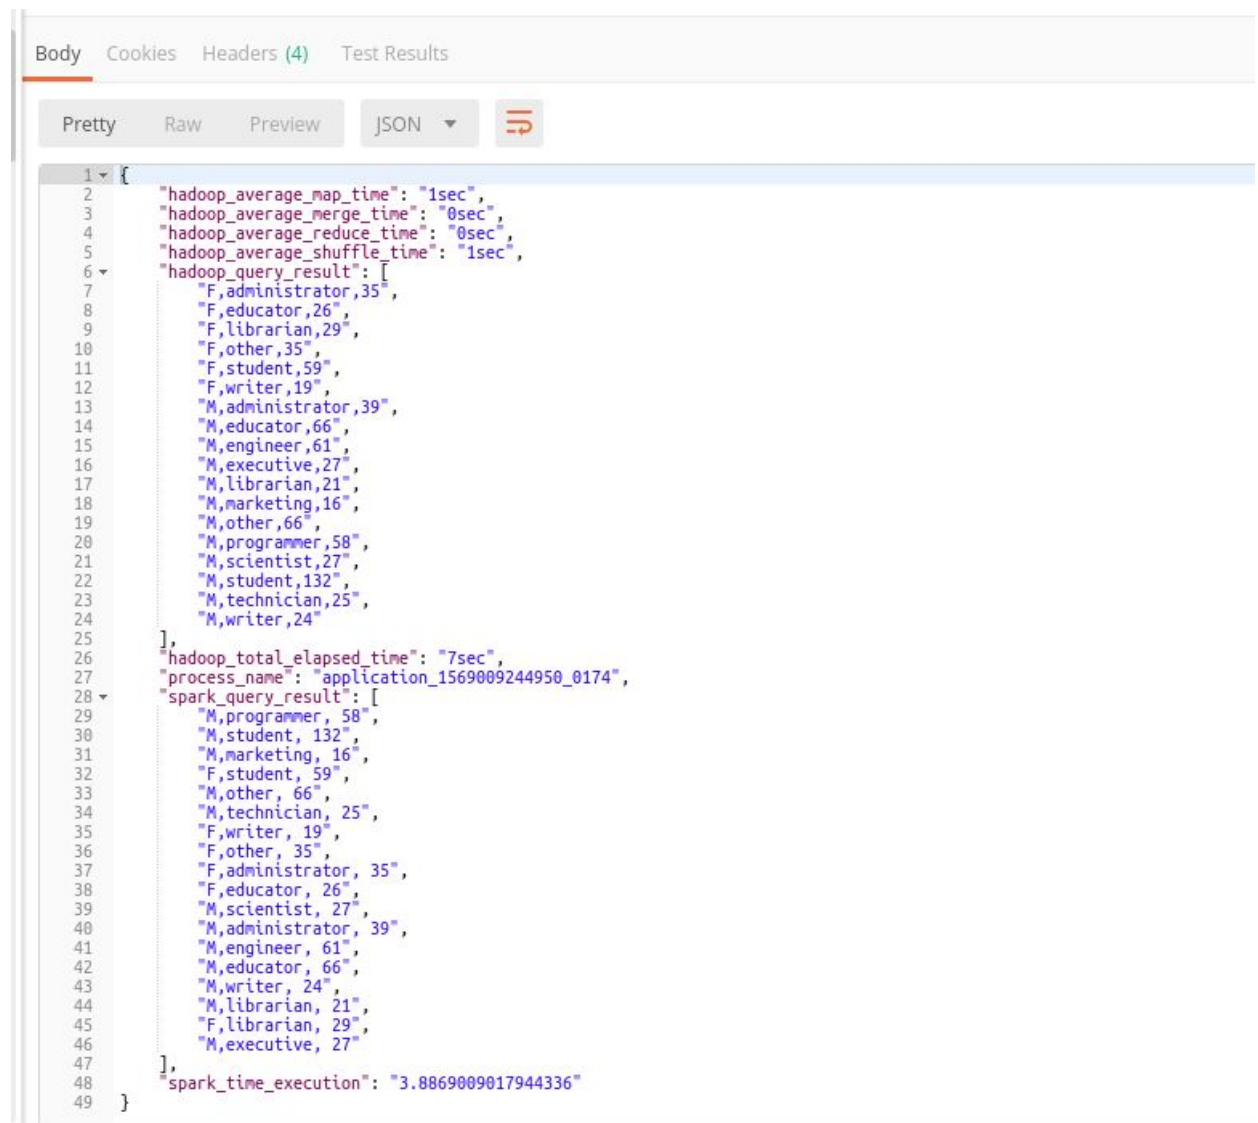


fig. Response from Server